

5

Unconstrained Optimization

In unconstrained optimization we want to find the global minimum of a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$. Without further assumptions, this problem is undecidable, so in order to be able to design (and analyze) algorithms, we need to restrict the class of functions we can tackle.

For starters, we make two basic assumptions:

1. the function f admits an optimal solution x^* of value p^* .
2. the gradient ∇f cannot change arbitrarily fast. More formally, we assume that ∇f is Lipschitz continuous, i.e.:

$$\|\nabla f(x) - \nabla f(y)\| \leq M\|x - y\| \quad (1)$$

for some fixed constant $M \geq 0$ in the domain of the function. For C^2 functions, this is equivalent to $\nabla^2 f(x) \preceq MI$.

Since the function is sufficiently smooth (at least C^1 given the assumptions above), the first order conditions for optimality require that $\nabla f(x^*) = 0$ on any optimal solution x^* . So, at least in principle, a first step is to solve such system of n nonlinear equations in n variables. Unlike linear systems, this in general cannot be solved analytically, so we resort to iterative methods that compute a sequence of points:

$$x^0, x^1, x^2, \dots$$

converging to a solution satisfying the first order condition to the limit:

$$\lim_{k \rightarrow +\infty} \nabla f(x^k) = 0$$

In the following, we will only consider iterative methods that maintain primal feasibility, i.e., $x^k \in \text{dom } f$ for all k . Also, we will only consider *descent* methods, i.e., iterative schemes where $f(x^{k+1}) < f(x^k)$. Note that, given a feasible starting point x^0 , we can without loss of generality restrict our search to the sublevel set $S = \{x \in \text{dom } f | f(x) \leq f(x^0)\}$: we will often assume that this set S is compact, i.e., closed and bounded, as this vastly increases the set of functions for which condition (1) is satisfied—it just need to hold for the points in S .

In the computer science meaning of the word: there is no algorithm with guaranteed termination.

This also means that f is bounded below by p^ .*

For convex functions, this is also a sufficient condition for global optimality!

This means that the objective function monotonically improves at every step.

5.1 Iterative methods

Without loss of generality, all iterative methods share the same pattern:

$$x^{k+1} = x^k + t^k \Delta x^k$$

This is indeed referred to as step length or step size.

In other words, the next point x^{k+1} in the sequence is obtained by the current point x^k and taking a step of length $t^k \geq 0$ along the *search direction* Δx^k . The algorithm will start from some initial point x^0 , and continue until some stopping criterion is met, usually based on the norm the gradient, e.g., $\|\nabla f(x^k)\| \leq \varepsilon$. A pseudocode is available in Algorithm 3.

Algorithm 3: Generic iterative scheme.

```

let  $x \in \text{dom } f$  be a starting point
while  $\|\nabla f(x)\| > \varepsilon$  do
    Determine direction  $\Delta x$ 
    Choose step size  $t$ 
    Update  $x \leftarrow x + t\Delta x$ 
end
return  $x$ 
```

Of course, different choices of the search direction and/or stepsize give rise to different algorithms, with different convergence properties.

5.1.1 *Choice of step length*

Suppose the search direction Δx is already chosen: how do we pick a step size? We have two conflicting goals here: on the one hand, the direction chosen (e.g., the gradient) might give a valid approximation of the function only locally, so we cannot take too large of a step; on the other hand, we also don't want the step to be too small when a larger one could have been beneficial, as that can slow down convergence significantly.

There are basically two options:

Remember, x and Δx are fixed.

1. *exact line search*: solves the one-dimensional optimization problem

$$t = \arg \min_{s \geq 0} f(x + s\Delta x)$$

This is used when t can be computed analytically and/or its costs is low compared to finding Δx .

2. *line search (inexact)*: does not look for the optimal t , just aims at reducing f enough.

A very simple, yet effective, strategy of line search is called *backtracking* line search, based on the Armijo-Goldstein conditions. It is parametrized by two constants, $0 < \alpha < \frac{1}{2}$ and $0 < \beta < 1$, and it consists in starting with a step length $t = 1$, and then reducing it geometrically, with a factor β , until the stopping condition $f(x + t\Delta x) \leq f(x) + \alpha t \nabla f(x)^\top \Delta x$ is met. A pseudocode is available in Algorithm 4.

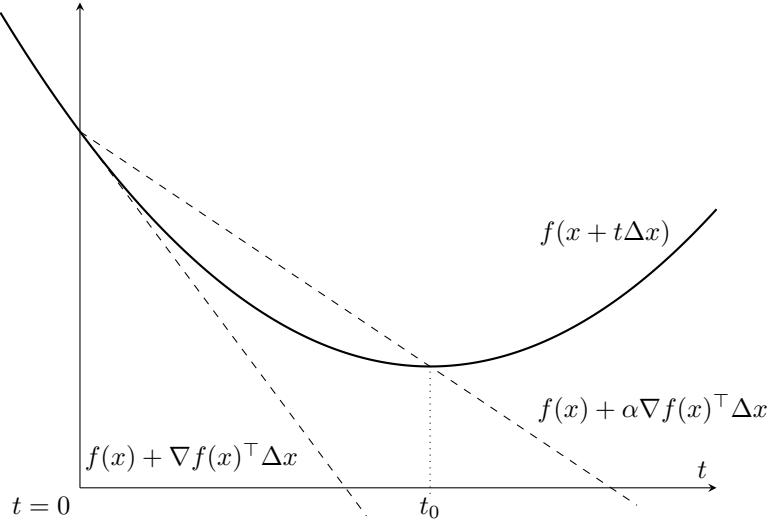


Figure 5.1: Backtracking line search.

Algorithm 4: Backtracking line search.

```

let  $t = 1$ 
while  $f(x + t\Delta x) > f(x) + \alpha t \nabla f(x)^\top \Delta x$  do
|  $t \leftarrow \beta t$ 
end
return  $t$ 
```

A graphical representation of the stopping condition is depicted in Figure 5.1. The idea is that we decrease the slope of the tangent at $t = 0$ by a factor α , and look for the largest value of t such that f is below the modified line. If t_0 is indeed the largest such value, it can be shown that when line search stops we either have $t = 1$ (the initial step already satisfies the condition) or $t \in (\beta t_0, t_0]$ (we are at most a factor of β away from t_0). Either way, we have that $t \geq \min\{1, \beta t_0\}$, so t cannot be too small.

5.2 Gradient Descent

How do we pick a search direction? If the function to optimize is convex, then by convexity alone we have:

$$f(x^{k+1}) \geq f(x^k) + \nabla f(x^k)^\top \underbrace{(x^{k+1} - x^k)}_{t^k \Delta x^k}$$

so if we want a descent direction it must hold that:

$$\nabla f(x^k)^\top \Delta x^k < 0 \tag{2}$$

The most obvious choice is $\Delta x = -\nabla f(x)$, so the negative gradient. It clearly satisfies (2), as we get

$$\nabla f(x^k)^\top \Delta x^k = -\|\nabla f(x^k)\|^2 < 0$$

More on this later.

unless the gradient is null, in which case we are done.

Even if the function is non-convex, choosing the negative gradient is a very natural choice, as the negative gradient is the direction of steepest descent w.r.t. the local first order approximation of the function.

So the question becomes: does such a method converge? If it does, at which speed? Note that, in all those schemes, the overall complexity is the product of two terms: the *iteration cost*, i.e., how expensive each individual iteration is, and the *iteration complexity*, i.e., how many iterations are needed in order to reach a desired accuracy.

5.2.1 Descent Lemma

The condition (1) alone implies the so-called *descent lemma*:

$$f(y) \leq f(x) + \nabla f(x)^\top (y - x) + \frac{M}{2} \|y - x\|^2 \quad (3)$$

Proof.

$$\begin{aligned} f(y) &= f(x) + \int_0^1 \nabla f(x + a(y - x))^\top (y - x) da && \text{(fund. thm. calculus)} \\ &= f(x) + \nabla f(x)^\top (y - x) + \int_0^1 (\nabla f(x + a(y - x)) - \nabla f(x))^\top (y - x) da && (\pm \text{constant}) \\ &\leq f(x) + \nabla f(x)^\top (y - x) + \int_0^1 \|\nabla f(x + a(y - x)) - \nabla f(x)\| \|y - x\| da && \text{(Cauchy-Schwarz)} \\ &\leq f(x) + \nabla f(x)^\top (y - x) + \int_0^1 M \|x + a(y - x) - x\| \|y - x\| da && \text{(Lipschitz)} \\ &= f(x) + \nabla f(x)^\top (y - x) + \int_0^1 aM \|y - x\|^2 da && \text{(factor out)} \\ &= f(x) + \nabla f(x)^\top (y - x) + \frac{1}{2} M \|y - x\|^2 \end{aligned}$$

□

The descent lemma implies that we have a quadratic upper bound on our function f , which is minimized by:

$$\hat{y} = x - \frac{1}{M} \nabla f(x)$$

A much simpler strategy than backtracking line search or exact methods.

This means that gradient descent with a constant step

$$t^k = \frac{1}{M}$$

will always achieve some improvement that is enough for convergence (albeit a slow one).

Proof. With a step length of $\frac{1}{M}$ and a direction of $-\nabla f(x^k)$, the difference $x^{k+1} - x^k$ is equal to $-\frac{1}{M} \nabla f(x^k)$. Plugging this into the inequality of the descent lemma we get:

$$\begin{aligned} f(x^{k+1}) &\leq f(x^k) + \nabla f(x^k)^\top (x^{k+1} - x^k) + \frac{M}{2} \|x^{k+1} - x^k\|^2 \\ &= f(x^k) - \frac{1}{M} \|\nabla f(x^k)\|^2 + \frac{M}{2} \frac{1}{M^2} \|\nabla f(x^k)\|^2 \\ &= f(x^k) - \frac{1}{2M} \|\nabla f(x^k)\|^2 \end{aligned}$$

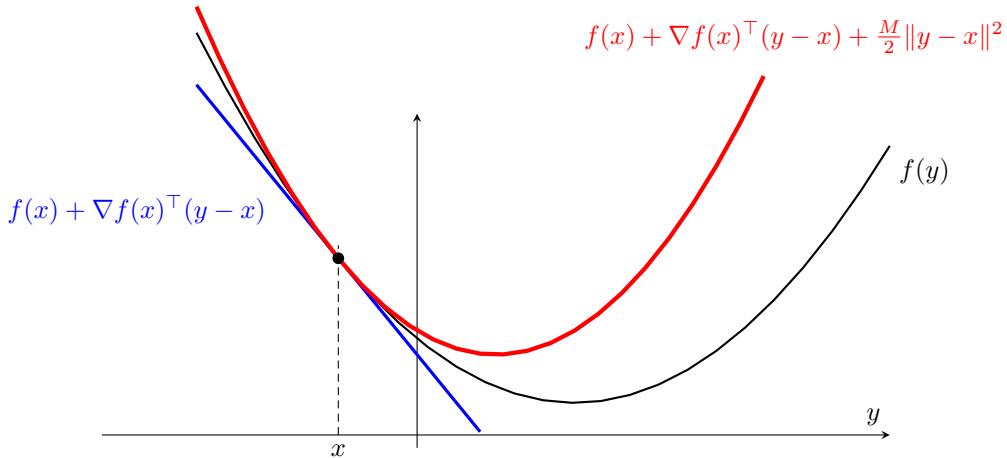


Figure 5.2: Descent lemma.

so we have a guaranteed progress of $\frac{1}{2M} \|\nabla f(x^k)\|^2$. This is enough to get our first convergence result.

Note that in practice we do not know M , but we can guess it. Start with $M = 1$ and double it until

$$\begin{aligned} f(x^{k+1}) &\leq f(x^k) - \frac{1}{2M} \|\nabla f(x^k)\|^2 \\ \|\nabla f(x^k)\|^2 &\leq 2M(f(x^k) - f(x^{k+1})) \end{aligned}$$

Now we can just average the last inequalities over the first s terms:

$$f\left(x^k - \frac{\nabla f(x^k)}{M}\right) \leq f(x^k) - \frac{\|\nabla f(x^k)\|^2}{2M}$$

In practice this is even better than knowing the real M , but the convergence is very slow. Backtracking line search is way better!

$$\underbrace{\frac{1}{s} \sum_{k=0}^{s-1} \|\nabla f(x^k)\|^2}_{\text{minimum gives lower bound}} \leq \underbrace{\frac{2M}{s} \sum_{k=0}^{s-1} (f(x^k) - f(x^{k+1}))}_{\text{telescoping sum}}$$

and obtain

$$\begin{aligned} \min_{k=0, \dots, s-1} \|\nabla f(x^k)\|^2 &\leq \frac{2M}{s} (f(x^0) - f(x^s)) \\ &\leq \frac{2M}{s} (f(x^0) - p^*) \end{aligned}$$

So in s iterations we find at least one iteration where

$$\|\nabla f\|^2 = O\left(\frac{1}{s}\right)$$

This is not necessarily the last iteration, but this is enough to show that

$$\liminf_{s \rightarrow +\infty} \|\nabla f\|^2 = 0$$

□ Note that we didn't even assume convexity! For non-convex functions, we will just converge to a stationary point.

What does this imply for accuracy? We have that:

$$\frac{2M(f(x^0) - p^*)}{s} \leq \varepsilon$$

and thus

$$s \geq \frac{2M(f(x^0) - p^*)}{\varepsilon}$$

so $s = O(\frac{1}{\varepsilon})$. This is a *bad* iteration complexity, as it is exponential in the number of digits of accuracy, i.e., $\log(\frac{1}{\varepsilon})$. Note that the convergence above is obtained with a constant step size, but it applies to exact line search as well, as that can only make a larger progress at each step.

5.2.2 Polyak-Lojasiewicz Inequality

The iteration complexity if $O(\frac{1}{\varepsilon})$ is the best we can prove without further assumptions on the function f we are optimizing. We can achieve a much better complexity if we assume the so-called Polyak-Lojasiewicz (PL) inequality:

$$\frac{1}{2} \|\nabla f(x)\|^2 \geq m(f(x) - p^*)$$

for all $x \in S$ and some $m > 0$.

For functions of class C^2 , strong convexity implies the PL inequality, with the same constant m :

$$\nabla^2 f(x) \succeq mI$$

Assuming the PL inequality, we can prove a much better iteration complexity for the same gradient scheme with constant step size $t^k = \frac{1}{M}$. Starting again from:

$$f(x^{k+1}) \leq f(x^k) - \frac{1}{2M} \|\nabla f(x^k)\|^2$$

we can now exploit that

$$-\|\nabla f(x^k)\|^2 \leq -2m(f(x^k) - p^*)$$

and thus, combining the two:

$$f(x^{k+1}) - p^* \leq f(x^k) - p^* - \frac{m}{M}(f(x^k) - p^*)$$

It is easy to show that $m \leq M$. which directly implies that at each iteration the error decreases at least by a factor of $(1 - \frac{m}{M}) < 1$. Iterating over s steps, we finally get:

$$f(x^s) - p^* \leq \left(1 - \frac{m}{M}\right)^s (f(x^0) - p^*)$$

so $f(x^s) - p^* = O(\rho^s)$ for some $\rho = (1 - m/M) < 1$. This is called exponential convergence, and it implies that to achieve an error of ε we need $O(\log \frac{1}{\varepsilon})$ iterations. Indeed, if we impose that the error after s iterations is bounded by ε and rearrange we get:

And linear in the number of digits of accuracy

Again, this applies also to exact line search and, with minimal adjustments, to backtracking line search as well.

$$s \geq \frac{\log \frac{(f(x^0) - p^*)}{\varepsilon}}{\log \frac{1}{\rho}} = O\left(\log \frac{1}{\varepsilon}\right) \quad (4)$$

Note that in (4), the numerator depends on how far from the optimal value the initial point x^0 is, relative to the desired tolerance ε , while the denominator depends on the constants m and M and thus, ultimately, on the structure of the function: if the so called condition number M/m is large, then ρ is very close to 1 and convergence is very slow. On the other hand, if the condition number is small, then ρ is close to zero and convergence is very fast.

Practical considerations. The analysis above is verified in practice in computational evaluations, and it turns out to be quite accurate. In particular:

- the effect of α and β within backtracking line search is noticeable but not dramatic;
- the same can be stated about the difference between exact line search and backtracking line search;
- convergence is truly sensitive to the condition number M/m , and it becomes painfully slow if this number is even modestly large (say, ≥ 1000).

So, we can conclude that the main advantage of gradient descent is that it is very simple and the cost of each iteration is small, but at the cost of having an iteration complexity which is too sensitive to the condition number, and potential very slow if this number is even modestly large.

Note that the iteration complexity does not depend directly on the dimension n , although the constants m and M might. Also, note that in practice m and M are unknown: they are just conceptual tools used in the convergence analysis, to understand the behaviour of the method.

5.3 Stochastic Gradient Descent

There are many applications where the objective function $f(x)$ to be minimized can be expressed as a sum (or average) of functions over a given set D :

$$f(x) = \frac{1}{|D|} \sum_{i \in D} f_i(x)$$

The most important example is machine learning (ML), where D is the set of samples/observations/data points. In this case, computing the gradient, while still technically linear in the dimension n , can be prohibitively expensive, as it also depends linearly on $d = |D|$, and we usually have $d \gg n$.

Its complexity is $O(nd)$.

In those cases, gradient descent is modified to replace the gradient with a cheaper approximation, that can be obtained by only using a (small) subset of the functions in the set D at each iteration (the so-called *sampled gradient*): this variant is called *stochastic gradient descent* (SGD), and it is one of the backbones of modern ML.

While quite natural and intuitive, the sample gradient is *not* a gradient of $f(x)$, so monotonic decrease is not guaranteed—and indeed does not happen in practice. Can we still prove convergence of the overall scheme? At which rate?

Our analysis is based on the simplest variant of SGD, where at each iteration we pick only a single sample from D . This is called *online SGD*. The key observation for proving convergence is that the sampled gradient is an *unbiased estimate* of the true gradient, assuming the sample is drawn *uniformly at random*. More formally:

$$\begin{aligned}
E[\nabla f_{i_k}(x)] &= \sum_{i \in D} p[i_k = i] \nabla f_i(x) \\
&= \frac{1}{|D|} \sum_{i \in D} \nabla f_i(x) \\
&= \nabla f(x)
\end{aligned}$$

However, we will see that in this case we will need the step size to converge to zero (there is no convergence with a constant step size).

Proof. Let's start again with the descent lemma:

$$f(x^{k+1}) \leq f(x^k) + \nabla f(x^k)^\top (x^{k+1} - x^k) + \frac{M}{2} \|x^{k+1} - x^k\|^2$$

and plug in $x^{k+1} - x^k = -t_k \nabla f_{i_k}(x^k)$. We obtain:

$$f(x^{k+1}) \leq f(x^k) - t_k \nabla f(x^k) \nabla f_{i_k}(x^k) + t_k^2 \frac{M}{2} \|\nabla f_{i_k}(x^k)\|^2$$

We can now take the expectation on both sides, assuming again uniform sampling:

$$\begin{aligned}
E[f(x^{k+1})] &\leq f(x^k) - t_k \nabla f(x^k) \underbrace{E[\nabla f_{i_k}(x^k)]}_{\nabla f(x^k)} + t_k^2 \frac{M}{2} E[\|\nabla f_{i_k}(x^k)\|^2] \\
&= f(x^k) - t_k \|\nabla f(x^k)\|^2 + t_k^2 \frac{M}{2} E[\|\nabla f_{i_k}(x^k)\|^2]
\end{aligned}$$

The last term is the expression above is *bad*:

- it reduces progress if the gradients are different;
- it makes it very risky to take large steps.

On the other hand, notice that if t_k is small, then $t_k \gg t_k^2$. For the sake of simplicity, let us make the *finite variance* assumption:

$$E[\|\nabla f_{i_k}(x^k)\|^2] \leq \sigma^2$$

for some constant σ^2 . Then the expression simplifies to

$$E[f(x^{k+1})] \leq f(x^k) - t_k \|\nabla f(x^k)\|^2 + t_k^2 \frac{M}{2} \sigma^2$$

and, rearranging:

$$t_k \|\nabla f(x^k)\|^2 \leq f(x^k) - E[f(x^{k+1})] + t_k^2 \frac{M}{2} \sigma^2$$

At this point we can sum up the first $s + 1$ terms, and exploit again the usual tricks (take expectation, bound by minimum, and telescoping sums):

$$\sum_{k=0}^s t_k \|\nabla f(x^k)\|^2 \leq \sum_{k=0}^s (f(x^k) - E[f(x^{k+1})]) + \frac{M}{2} \sigma^2 \sum_{k=0}^s t_k^2$$

$$\sum_{k=0}^s t_k \underbrace{E[\|\nabla f(x^k)\|^2]}_{\text{lower bounded by min}} \leq \sum_{k=0}^s (\underbrace{E[f(x^k)] - E[f(x^{k+1})]}_{\text{telescoping sum}}) + \frac{M}{2}\sigma^2 \sum_{k=0}^s t_k^2$$

$$\min_{k=0,\dots,s} E[\|\nabla f(x^k)\|^2] \sum_{k=0}^s t_k \leq f(x^0) - \underbrace{E[f(x^s)]}_{\geq p^*} + \frac{M}{2}\sigma^2 \sum_{k=0}^s t_k^2$$

Finally, dividing by $\sum_{k=0}^s t_k$, we get

$$\min_{k=0,\dots,s} E[\|\nabla f(x^k)\|^2] \leq \frac{f(x^0) - p^*}{\sum_{k=0}^s t_k} + \frac{M}{2}\sigma^2 \frac{\sum_{k=0}^s t_k^2}{\sum_{k=0}^s t_k}$$

So, in order for the (expected) gradient to converge to zero, we need the two conditions:

$$\sum_k t_k \rightarrow +\infty \quad \sum_k t_k^2 \rightarrow 0$$

□

For example, with $t_k = a/k$ for some constant a , we have that

$$\sum_{k=1}^s t_k = O(\log s) \quad \sum_{k=1}^s t_k^2 = O(1)$$

so the error after s iterations is $O(1/\log s)$, which is worse than standard gradient descent. If, instead, we pick a constant step size $t_k = a$, we get an error of $O(1/s) + O(a)$, that does *not* converge to zero as $k \rightarrow +\infty$.

For strongly convex functions, a similar analysis gives, with the decreasing step size rule $t_k = 1/mk$, an error of $O(1/s)$ after s iterations, while with a constant step size $t_k = a < 1/2m$ we get that an error bound of $O(\rho^s) + O(a)$, which means linear convergence but only up to some accuracy proportional to a . Note that for both strongly convex and general functions, using a constant step size gives fast convergence at the beginning, but eventually the second (constant) term dominates: this justifies the heuristic used in practice of “*halving the step size if the process is not making progress*”, as the second term depends on a .

Again, worse than in the standard case.

5.3.1 Mini-batch SGD

As we have seen, standard gradient descent achieves faster convergence than SGD, but at the cost of d gradient evaluations per iteration, instead of just 1. Is there a middle ground between the two? A very natural approach consists in using more samples at each iteration, in order to have a more accurate approximation of the true gradient at x^k : this is called the *mini-batch* approach, and it consists in randomly sampling a subset B^k of indices at each iteration.

The k -th iteration of the method can thus be written as:

$$x^{k+1} = x^k - t_k \frac{1}{|B^k|} \sum_{i \in B^k} \nabla f_i(x^k)$$

Also called deterministic.

On modern parallel architectures, if the size of B^k is appropriately chosen, we can evaluate those gradients equally fast, so the iteration cost is not even higher.

Let us consider the mini-batch direction as the true gradient $\nabla f(x^k)$ plus some noise term e^k that depends on the sample we have chosen. By using a

constant step size $t_k = 1/M$, and repeating the steps of the descent lemma, we get the inequality:

$$f(x^{k+1}) \leq f(x^k) - \underbrace{\frac{1}{2M} \|\nabla f(x^k)\|^2}_{\text{good term}} + \underbrace{\frac{1}{2M} \|e^k\|^2}_{\text{bad term}}$$

As in the online case, the error term is usually the limiting factor for convergence. How does $|B^k|$ affect $\|e^k\|^2$? Under the usual finite variance assumption, it is easy to show that if we are sampling *with replacement*, we have:

$$E[\|e^k\|^2] = \frac{1}{|B^k|} \sigma^2$$

and thus, unsurprisingly, the error gets smaller as we increase the size of the mini-batch. Even better, if we sample *without replacement* from the full set of d gradients, we obtain:

$$E[\|e^k\|^2] = \frac{d - |B^k|}{d} \frac{1}{|B^k|} \sigma^2$$

and this goes to zero as $|B^k| \rightarrow d$. In other words, if we increase the size of the mini-batch over time at the appropriate rate, then can reduce the error $\|e^k\|^2$ fast enough that it is no longer the limiting factor, and we recover the convergence of standard gradient descent. For example, with the rule $|B^{k+1}| = |B^k|/\rho$ we can achieve $O(\rho^s)$ convergence, while with $|B^{k+1}| = |B^k| + a$ we match $O(1/s)$. Note, however, that those update rules are not used in implementations, as they increase the mini-batch size too quickly: in practice we prefer to keep the mini-batch size fixed (and carefully chosen to match the available hardware), and double its size only when stalling is detected.

5.4 Coordinate Descent

Stochastic gradient descent improves the iteration cost over gradient descent when the $|D|$ is large, but what about the case in which n is itself large? A very natural approach in this case is to do *coordinate descent*, where at each iteration k we only consider a single variable x_{j_k} :

$$x_{j_k}^{k+1} = x_{j_k}^k - t_k \nabla_{j_k} f(x^k)$$

while the remaining ones keep the same value, see Figure 5.3.

In general, this is a *provably bad* idea:

- it has a worst iteration complexity (as in SGD, the resulting direction is *not* a gradient);
- the iteration cost is the same, under the assumption that we still need to compute the full gradient before choosing a coordinate.

The picture changes with so-called *coordinate-friendly* functions, where evaluating a single component of the gradient is significantly cheaper than evaluating the full gradient. The simplest example is when a function is

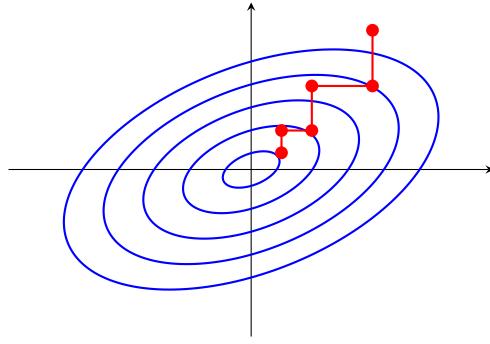


Figure 5.3: Coordinate descent.

separable, i.e., $f(x) = \sum_{j=1}^n f_j(x_j)$: here the full gradient costs $O(n)$, while a single component is only $O(1)$. A more interesting case is quadratic functions: the full gradient can be computed in $O(n^2)$, but a single component only requires $O(n)$, again a speedup of a factor of n .

What about convergence? Assume each $\nabla_j f(x)$ is M -Lipschitz:

This is indeed a very special case: if the function is fully separable, you would decompose it into independent pieces and do coordinate descent anyway.

For \mathbb{C}^2 functions, this means $|\nabla_{jj}^2 f(x)| \leq M$.

for any $x \in S$ and for any j . This is a not stronger assumption than the one we usually make. We can consider the coordinate-wise equivalent of the descent lemma:

$$f(x^{k+1}) \leq f(x^k) + \nabla_{j_k} f(x^k)^\top (x^{k+1} - x^k) + \frac{M}{2} (x^{k+1} - x^k)^2$$

Note that x^{k+1} and x^k only differ in coordinate j_k . Plugging as usual:

$$x^{k+1} - x^k = -t_k \nabla_{j_k} f(x^k) e_{j_k}$$

and using a step size $t_k = 1/M$, we get:

$$f(x^{k+1}) \leq f(x^k) - \frac{1}{2M} |\nabla_{j_k} f(x^k)|^2$$

Let us assume again that the coordinate j_k is chosen uniformly at random at each iteration, and take the expectation:

$$\begin{aligned} E[f(x^{k+1})] &\leq E[f(x^k) - \frac{1}{2M} |\nabla_{j_k} f(x^k)|^2] \\ &= f(x^k) - \frac{1}{2M} \sum_{j=1}^n \frac{1}{n} |\nabla_{j_k} f(x^k)|^2 \\ &= f(x^k) - \frac{1}{2nM} \sum_{j=1}^n |\nabla_{j_k} f(x^k)|^2 \\ &= f(x^k) - \frac{1}{2nM} \|\nabla f(x^k)\|^2 \end{aligned}$$

So the convergence rate is n times *slower* than with standard gradient descent. If we can assume the PL inequality, we also get:

$$E[f(x^k)] - p^* \leq \left(1 - \frac{m}{nM}\right)^k (f(x^0) - p^*)$$

again, n times slower. In other words, we have n times cheaper iterations (in the best case), but n times more iterations: why bother? The trick is that the involved Lipschitz constants M are much smaller, so we can take longer steps. Note that this is with a constant step size, with backtracking line search the comparison is not straightforward. The same applies for other potential improvements: for a single coordinate we can usually afford more expensive methods (like, e.g., a Newton-like step), but this doesn't improve the theoretical convergence rate.

So far, we have assumed a random selection strategy: are there other options that give provably faster convergence? Here is a list of the common approaches:

- *cyclic selection*. It seems to work better in practice, but it is worse in theory.
- *random shuffling*. The idea is to choose a random permutation of the coordinates, and follow it for one cycle, then repeat (each time with a different random permutation). This seems to be faster in practice, but it is not fully understood in theory.
- *greedy*. This is also called the Gauss-Southwell rule, and consists in picking the coordinate with the largest coefficient in the gradient:

$$j_k = \arg \max_j \{|\nabla_j f(x^k)|\}$$

Note that, if implemented naïvely, this destroys the speed advantage on a single iteration, as we need to compute the full gradient to find the index giving the maximum value. This approach is thus viable only if, exploiting proper data structures and incremental updates, we can keep track of the maximum with the same complexity (usually, up to some logarithmic term). If this can be achieved, then we can indeed obtain a guaranteed progress of:

$$f(x^{k+1}) \leq f(x^k) - \frac{1}{2M} \|\nabla f(x^k)\|_\infty$$

Just with a different norm.

which is similar to the rate of standard gradient descent, so provably faster than random.

- *non-uniform sampling*. Another approach is to maintain a separate Lipschitz approximation M_j for each coordinate, and sample more often coordinates with larger values:

$$p[j_k = j] = \frac{M_j}{\sum_{i=1}^n M_i}$$

In this case we get slightly better constants in the analysis, but the overall behaviour is similar.

5.5 Newton's method

In this section we introduce a different approach for choosing a descent direction, that will lead to a method with a significantly faster convergence rate.

5.5.1 Steepest descent

The first-order approximation of a function f at x is given by:

$$f(x + v) \approx f(x) + \nabla f(x)^\top v$$

where the second term is the directional derivative of f at x in direction v . Since we are minimizing, we would like to pick a direction v such that the corresponding directional derivative is as small as possible. However, since v can always be arbitrarily scaled, we need to normalize the norm of the derivative w.r.t. the norm of v . For a given norm $\|\cdot\|$, we thus obtain the *steepest descent* direction at x as:

$$\Delta x_{nsd} = \arg \min_v \{\nabla f(x)^\top v \mid \|v\| = 1\}$$

For technical reasons, it is often convenient to scale the steepest descent direction using the dual norm $\|\cdot\|_*$, and obtain the *unnormalized* steepest descent direction:

$$\Delta x_{sd} = \|\nabla f(x)\|_* \Delta x_{nsd}$$

Given a norm $\|\cdot\|$ in \mathbb{R}^n , the dual norm is defined as

$$\|z\|_* = \sup\{z^\top x \mid \|x\| \leq 1\}$$

The steepest descent direction (normalized or unnormalized) thus depends on which norm we choose. If we choose the Euclidian norm, we obtain $\Delta x_{sd} = -\nabla f(x)$, so standard gradient descent is a steepest descent method with $\|\cdot\| = \|\cdot\|_2$. The Euclidian norm can be generalized to a much larger family, those of *quadratic norms*. Given a positive definite matrix $P \in S_{++}^n$, let us define $\|\cdot\|_P$ as:

$$\|z\|_P = (z^\top P z)^{\frac{1}{2}} = \|P^{\frac{1}{2}} z\|_2$$

In this case, it is easy to show that the corresponding unnormalized steepest descent direction is:

$$\Delta x_{sd} = -P^{-1} \nabla f(x)$$

We get the standard Euclidian norm with $P = I$.

This can be interpreted as a standard gradient descent after the change of coordinates $\bar{x} = P^{\frac{1}{2}} x$. While this change of coordinates can positively affect the condition number of the Hessian, and thus reduce the practical iteration complexity, if the norm is chosen once and for all, we still get *only* linear convergence.

5.5.2 Newton step

The Newton step (or direction) is defined as:

$$\Delta x_{nt} = -\nabla^2 f(x)^{-1} \nabla f(x)$$

If the function f is strongly convex, the Hessian $\nabla^2 f(x)$ is positive definite for any $x \in S$, and this is thus always a descent direction. This choice of the search direction can be justified in many different ways:

- Minimizer of second order approximation. The second-order approximation of f at x is given by:

$$f(x + v) \approx f(x) + \nabla f(x)^\top v + \frac{1}{2} v^\top \nabla^2 f(x) v$$

and this is a convex quadratic function of v , minimized by $v = \Delta x_{nt}$.

- Steepest direction in Hessian norm. $\Delta x_{nt} = \Delta x_{sd}$ for $P = \nabla^2 f(x)$.
- Linearized optimality conditions. The optimality conditions $\nabla f(x) = 0$ can be linearized near x to obtain:

$$\nabla f(x + v) \approx \nabla f(x) + \nabla^2 f(x)v = 0$$

which is a linear system with solution $v = \Delta x_{nt}$.

Note that even though we can interpret the Newton step as the steepest descent direction in Hessian norm, the key difference is that the chosen norm changes (in general) at every iteration.

5.5.3 *Affine invariance*

An important property of the Newton step is that it is *affine invariant*, i.e., the direction is *independent* of affine changes of coordinates. In other words, let us consider a function $f(x)$, a linear transformation $x = Ay$ encoded by the nonsingular square matrix $A \in \mathbb{R}^{n \times n}$, and their combination $g(y) = f(Ay) = f(x)$. The Newton step computed for g at y is:

$$\begin{aligned}\Delta y_{nt} &= -\nabla^2 g(y)^{-1} \nabla g(y) \\ &= -(A^\top \nabla^2 f(Ay) A)^{-1} A^\top \nabla f(Ay) \\ &= -A^{-1} \nabla^2 f(Ay)^{-1} \nabla f(Ay) \\ &= -A^{-1} \nabla^2 f(x)^{-1} \nabla f(x) \\ &= A^{-1} \Delta x_{nt}\end{aligned}$$

This means that a descent method based on the Newton step will always follow the same path, independent of affine transformations.

Note that gradient descent is not affine invariant.

5.5.4 *Convergence analysis*

Let us assume, as usual, that f is strongly convex, i.e.,

$$mI \preceq \nabla^2 f(x) \preceq MI$$

for any $x \in S$ and for some $0 < m \leq M$. In addition, we will assume that the Hessian itself is Lipschitz continuous on S :

$$\|\nabla^2 f(x) - \nabla^2 f(y)\| \leq L \|x - y\|$$

for some $L \geq 0$. Note that we trivially have $L = 0$ for quadratic functions. The main ingredient of the convergence analysis is given in the following lemma, that we state without proof:

Lemma 5.1. *Assume we do a descent method always choosing the Newton step and with backtracking line search. Then there exists two numbers $0 < \eta \leq \frac{m^2}{L}$ and $\gamma > 0$ such that:*

- if $\|\nabla f(x^k)\| \geq \eta$ then $f(x^{k+1}) \leq f(x^k) - \gamma$, i.e., the objective function improves by at least γ ;
- if $\|\nabla f(x^k)\| < \eta$ then backtracking line search will select $t_k = 1$ and we have

$$\frac{L}{2m^2} \|\nabla f(x^{k+1})\| \leq \left(\frac{L}{2m^2} \|\nabla f(x^k)\| \right)^2$$

Note that once the second condition is satisfied, it will always be true:

$$\begin{aligned} \|\nabla f(x^{k+1})\| &\leq \frac{L}{2m^2} \|\nabla f(x^k)\|^2 \\ &\leq \frac{L}{2m^2} \eta^2 \\ &\leq \frac{L}{2m^2} \frac{m^2}{L} \eta < \eta \end{aligned}$$

This means that once it becomes true at some iteration k it will stay true for all iterations $l \geq k$, and we will always take the full Newton step ($t = 1$). Applying the second inequality recursively we have:

$$\begin{aligned} \frac{L}{2m^2} \|\nabla f(x^l)\| &\leq \left(\frac{L}{2m^2} \|\nabla f(x^k)\| \right)^{2^{l-k}} \\ &\leq \left(\frac{L}{2m^2} \eta \right)^{2^{l-k}} \\ &\leq \left(\frac{1}{2} \right)^{2^{l-k}} \end{aligned}$$

Combining the above with the PL inequality we obtain:

$$f(x^l) - p^* \leq \frac{1}{2m} \|\nabla f(x^l)\|^2 \leq \frac{2m^3}{L} \left(\frac{1}{2} \right)^{2^{l-k+1}}$$

Once we are in this phase, convergence is thus extremely fast (we call it *quadratic convergence*). What about the first phase? Clearly the number of iterations there is at most:

$$\frac{f(x^0) - p^*}{\gamma}$$

So in total we have at most:

$$\frac{f(x^0) - p^*}{\gamma} + \log_2 \log_2 \left(\frac{2m^3}{L} \frac{1}{\varepsilon} \right)$$

For all practical purposes the second term is very small (e.g., < 10), so convergence is entirely dominated by the first phase: once we are relatively

close to an optimal solution, the second phase kicks in and it takes a few iterations to converge.

The analysis above is confirmed by empirical evaluations: the number of iterations required to converge is relatively small in practice (in the order of some hundreds, at most), and mostly unaffected by the dimension n and the type of line search. In other words, the iteration complexity of the Newton's method is significantly better than that of gradient descent. On the other hand, the iteration cost is much higher: we need to compute the Hessian $H = \nabla^2 f(x)$ and solve the linear system $Hv = -\nabla f(x)$, which has a cost both in terms of memory and time which is in general at least $O(n^2)$, and thus completely unacceptable for some large applications.

Finally, we note that the analysis given above is not very satisfactory: it depends on quantities we do not know in practice (m, M, L), and that are not even affine invariant, while the overall is. However, there are more refined analyses, based on the theory of self concordant functions,, that give bounds that only depend on known parameters $(\alpha, \beta, \varepsilon)$ and that are affine invariant.

Functions where $|f'''(x)| \leq 2f''(x)^{3/2}$.